

PEA Data Documentation

Purpose and sources

PEA assesses whether supplied evidence supports a performance claim. Its data are synthetic development and evaluation material, not employee records or fine-tuning data. Source identifiers below are relative to the submission root. The release manifest identifies copied files and their hashes.

Material	Recorded size	Source and purpose
Current development set	40 cases, 10 per class	pea-backend/data/dev; AI-draft development labels
PEA-60	60 cases, 15 per class	pea-backend/data/holdout; formal ABCD comparison
Agent scenarios	12	pea-backend/data/agent_eval; functional retrieval scenarios
Mock goal updates	12 rows	pea-backend/data/evidence/goal_updates.jsonl
Mock feedback	4 rows	pea-backend/data/evidence/feedback.jsonl
Mock one-to-one notes	4 rows	pea-backend/data/evidence/one_to_one.jsonl
Reviewed PEA-20	20 cases	reference_materials/PEA20_Human_Reviewed_Anchor_Set_English_Submission.xlsx

The reviewed PEA-20 workbook has Reviewed_Cases, Model_Input and Audit_Trail sheets. It is distinct from the 20 template seed IDs in pea-backend/data/dev/seed_case_ids.json. The newly located workbook updates the older evidence summary, which had not found it. Do not treat template seeds as recovered human-reviewed anchors.

Current dev40 differs from historical dev140. The saved report

pea-backend/results/experiment_71d92befcef44aea975f85e13820c302/report.json contains the historical 280 method-case results. It must not be merged with current inputs or PEA-60. The previous draft holdout is separately retained under data/backups/holdout-before-pea60.

Human review and generation

The PEA-60 dataset_manifest.json records 33 explicit labels and 27 blank cells interpreted as agreement with AI drafts under the user's instruction, with zero final label changes. Describe this as human-reviewed AI-drafted labels, not independent human-authored Gold. There is no recovered inter-rater reliability measurement.

Both candidate and reviewed workbooks were located in Downloads and copied under reference_materials. Their states differ: the English candidate's Human_Review sheet can contain blank labels and Not Reviewed status. It is not interchangeable with the reviewed Chinese workbook. The manifest's 33/27 process is disclosed as recorded provenance; a complete cell-by-cell reconstruction was not performed for this accelerated delivery.

The holdout manifest attributes generation to GPT-5.6 Sol (OpenAI). Original generation prompts and invocation records were not recovered, so this is a supplied attribution rather than independent verification of the precise generator. The general data/manifest.json describes deterministic development templates; it must not replace holdout-specific provenance. These datasets support inference evaluation, not model training.

Schema, transformations and separation

Files use UTF-8 JSONL with separate inputs.jsonl and labels.jsonl. Current dev and Agent directories also contain annotations.jsonl. PEA-60's performance criterion and definition are adapted by app/dataset_inputs.py. app/schemas.py validates fields, lengths, enums and dates. Case IDs identify cases; employee IDs and review periods constrain retrieval; reference standards supply comparison context. Evidence identifiers support citations and deduplication. Original evidence IDs differ from runtime citation IDs.

app/models.py projects only approved case facts to the model. Gold labels, rationales, boundary types and human-review metadata are excluded. Neutral case/employee identifiers are backend scope data, not model-visible answers. app/evaluation.py persists predictions before loading labels. Hashing Gold bytes for integrity does not mean placing Gold in model context. Prior selected-case smoke exposure must still be disclosed.

The local evidence library has authorization, employee/case, source and date filters followed by bilingual keyword matching. It is not a live HRIS or vector database. Duplicate, contradictory, wrong-scope and no-result fixtures test defined behaviors; they do not establish general retrieval recall. Missing metadata is not invented.

Validation and limits

The new verification/data_inventory.json records direct JSONL counts and language fields. Full schema validation, installation and application tests were NOT RUN during the accelerated delivery. The existing command is python data/validate_data.py from pea-backend after dependency installation. Complete safe JSONL inputs and labels are explicitly included, independent of .gitignore.

No real employee data or production credentials are included. A redistribution licence for every supplied workbook was not established; confirm rights before public release. Synthetic data, AI-draft labels, a small sample and limited review evidence restrict generalization. Missing generation artifacts must remain gaps, not reconstructed originals.